

Kubernetes Essentials Tutorial

Hello Minikube - Your First Cluster

What is Minikube? Minikube creates a local Kubernetes cluster for development and learning.

Setup & Start:

```
# Start minikube
```

```
minikube start
```

```
# Open the Kubernetes dashboard
```

```
minikube dashboard
```

Create Your First Deployment:

```
# Create a deployment
```

```
kubectl create deployment hello-node --image=registry.k8s.io/e2e-test-images/agnhost:2.53 -- /agnhost netexec --http-port=8080
```

```
# View deployment
```

```
kubectl get deployments
```

```
# View pods
```

```
kubectl get pods
```

Expose Your App:

```
# Create a service to expose the deployment
```

```
kubectl expose deployment hello-node --type=LoadBalancer --port=8080
```

```
# Access the service
```

```
minikube service hello-node
```

Clean up:

```
kubectl delete service hello-node
```

```
kubectl delete deployment hello-node
```

Deploying Applications with kubectl

Deployments manage Pods and ensure your application stays running. They handle scaling, rolling updates, and self-healing.

Basic kubectl Commands:

```
# Create deployment
```

```
kubectl create deployment kubernetes-bootcamp --image=gcr.io/google-samples/kubernetes-bootcamp:v1
```

```
# View deployments
```

```
kubectl get deployments
```

```
# View pods
```

```
kubectl get pods
```

```
# Create proxy to access pods
```

```
kubectl proxy
```

Access your application:

```
# Get pod name
export POD_NAME=$(kubectl get pods -o go-template --template
'{{range .items}}{{.metadata.name}}\n{{end}}')
```

```
# Access through proxy
```

```
curl http://localhost:8001/api/v1/namespaces/default/pods/$POD_NAME:8080/proxy/
```

Exploring Pods and Nodes

Pods are the smallest deployable units in Kubernetes, containing one or more containers.

Nodes are worker machines (virtual or physical) that run your Pods.

Key Commands:

```
# Get detailed pod information
```

```
kubectl describe pods
```

```
# View logs
```

```
kubectl logs <pod-name>
```

```
# Execute commands in pod
```

```
kubectl exec <pod-name> -- env
```

```
kubectl exec -ti <pod-name> -- bash
```

Exposing Applications with Services

Services provide stable network access to Pods, handling load balancing and service discovery.

Service Types:

- **ClusterIP**: Internal cluster access only (default)
- **NodePort**: Exposes service on each node's IP at a static port
- **LoadBalancer**: Exposes service externally using cloud provider's load balancer

Create and expose a service:

```
# Expose deployment as NodePort service
```

```
kubectl expose deployment/kubernetes-bootcamp --type="NodePort" --port=8080
```

```
# Get service details
```

```
kubectl describe services/kubernetes-bootcamp
```

```
# Get the node port
```

```
export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
```

```
# Access the service
```

```
curl http://$(minikube ip):$NODE_PORT
```

Working with Labels:

```
# Label a pod
```

```
kubectl label pods <pod-name> version=v1
```

```

# Query by label
kubectl get pods -l app=kubernetes-bootcamp
kubectl get services -l app=kubernetes-bootcamp

# Delete service by label
kubectl delete service -l app=kubernetes-bootcamp

Scaling Applications
Manual Scaling:
# Scale to 4 replicas
kubectl scale deployments/kubernetes-bootcamp --replicas=4

# Check deployment status
kubectl get deployments

# View replica sets
kubectl get rs

# Check load balancing (run multiple times)
curl http://$(minikube ip):$NODE_PORT

Scale down:
# Reduce to 2 replicas
kubectl scale deployments/kubernetes-bootcamp --replicas=2

Rolling Updates
Update application version:
# Update to version 2
kubectl set image deployments/kubernetes-bootcamp \
  kubernetes-bootcamp=docker.io/jocatalin/kubernetes-bootcamp:v2 \
  kubernetes- \
  bootcamp=v2

# Check rollout status
kubectl rollout status deployments/kubernetes-bootcamp

# Verify update
curl http://$(minikube ip):$NODE_PORT

Rollback if needed:
# Rollback to previous version
kubectl rollout undo deployments/kubernetes-bootcamp

# Check pod status after rollback
kubectl get pods

Essential kubectl Commands Summary
# Cluster info
kubectl version
kubectl get nodes

```

```
# Deployments
kubectl create deployment <name> --image=<image>
kubectl get deployments
kubectl describe deployment <name>
kubectl delete deployment <name>

# Pods
kubectl get pods
kubectl describe pods <name>
kubectl logs <name>
kubectl exec -ti <name> -- bash

# Services
kubectl expose deployment <name> --type=<type> --port=<port>
kubectl get services
kubectl describe service <name>

# Scaling
kubectl scale deployments/<name> --replicas=<number>

# Updates
kubectl set image deployments/<name> <container>=<image>:<tag>
kubectl rollout status deployments/<name>
kubectl rollout undo deployments/<name>
```

Open the Dashboard

Open the Kubernetes dashboard. You can do this two different ways:

Launch a browser

[URL copy and paste](#)

Open a **new** terminal, and run:

```
# Start a new terminal, and leave this running.  
minikube dashboard
```

Now, switch back to the terminal where you ran `minikube start`.

Note:

The `dashboard` command enables the dashboard add-on and opens the proxy in the default web browser. You can create Kubernetes resources on the dashboard such as Deployment and Service.

To find out how to avoid directly invoking the browser from the terminal and get a URL for the web dashboard, see the "URL copy and paste" tab.

By default, the dashboard is only accessible from within the internal Kubernetes virtual network. The `dashboard` command creates a temporary proxy to make the dashboard accessible from outside the Kubernetes virtual network.

To stop the proxy, run `ctrl+c` to exit the process. After the command exits, the dashboard remains running in the Kubernetes cluster. You can run the `dashboard` command again to create another proxy to access the dashboard.

Using kubectl to Create a Deployment

Objectives

- Learn about application Deployments.
- Deploy your first app on Kubernetes with kubectl.

Kubernetes Deployments

A Deployment is responsible for creating and updating instances of your application.

Note:

This tutorial uses a container that requires the AMD64 architecture. If you are using minikube on a computer with a different CPU architecture, you could try using minikube with a driver that can emulate AMD64. For example, the Docker Desktop driver can do this.

Once you have a [running Kubernetes cluster](#), you can deploy your containerized applications on top of it. To do so, you create a Kubernetes **Deployment**. The Deployment instructs Kubernetes how to create and update instances of your application. Once you've created a Deployment, the Kubernetes control plane schedules the application instances included in that Deployment to run on individual Nodes in the cluster.

Viewing Pods and Nodes

Objectives

- Learn about Kubernetes Pods.
- Learn about Kubernetes Nodes.
- Troubleshoot deployed applications.

Kubernetes Pods

A Pod is a group of one or more application containers (such as Docker) and includes shared storage (volumes), IP address and information about how to run them.

When you created a Deployment in [Module 2](#), Kubernetes created a **Pod** to host your application instance. A Pod is a Kubernetes abstraction that represents a group of one or more application containers (such as Docker), and some shared resources for those containers. Those resources include:

- Shared storage, as Volumes
- Networking, as a unique cluster IP address
- Information about how to run each container, such as the container image version or specific ports to use